Zachary Werder

A2 Write Up

1. I decided to use the prior-value context-adaptive source model for the video data.

2. I thought this scheme would be good because prior context seems very applicable for pixel values. There might be some consistency between pixels across frames. It seemed like an easier to adapt model for film data.

3. Most of the alterations I tried ended up performing terribly, and I ended up pretty much only saving between the text compression what was afforded by knowing the exact specifications of the video. I was looking to achieve compression down to 50%, and got down to around 70%. So all and all not the best.

4. As expected, the context- adaptive scheme seemed to perform best on the data. I would anticipate my scheme working better than text encoders, my thinking is video is just easier to compress.

5. I think changes I could make have to do with allowing an encoder that takes advantage of the perceptibility of form to create a more functional lossy image. The human eye would not notice small differences here and there in color. But then again that's assuming the compression is meant for people and not some sort of AI. I would think the context of the application of the encoder is important, and flexibility between different contexts a sign of greater success.

   Film is pretty dear to me, and finding ways to compress film for future generations is important. I am curious how one might go about doing so in none digital applications, if compression in that context even makes sense. If our digital world crumbles what will be left?